

Experiment - 6 Naive Bayes Classification

Code

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from sklearn.preprocessing import LabelEncoder
from sklearn.datasets import load_iris

def implement_naive_bayes(dataset_name):
    print(f"\n{' '*15} Naive Bayes on {dataset_name} Dataset {' '*15}")

    # Load the chosen dataset
    if dataset_name == "Iris":
        data = load_iris()
        X = data.data
        y = data.target
        feature_names = data.feature_names
        target_names = data.target_names
    elif dataset_name == "Wine":
        data = load_wine()
        X = data.data
        y = data.target
        feature_names = data.feature_names
        target_names = data.target_names

    # Split data into training and test sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)

    # Display dataset information
    print(f"\nDataset Information:")
    print(f"Number of samples: {X.shape[0]}")
    print(f"Number of features: {X.shape[1]}")
    print(f"Number of classes: {len(target_names)}")
    print(f"Classes: {target_names}")
    print(f"Features: {feature_names}")
    print(f"\nTraining set size: {X_train.shape[0]}")
    print(f"Test set size: {X_test.shape[0]}")

    # Initialize and train the Naive Bayes classifier
    print("\nTraining Naive Bayes classifier...")
    naive_bayes = GaussianNB()
    naive_bayes.fit(X_train, y_train)
```

```

# Make predictions
y_pred = naive_bayes.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"\nAccuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")

# Display confusion matrix
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)

# Display classification report
print("\nClassification Report:")
report = classification_report(y_test, y_pred, target_names=target_names)
print(report)

# Visualize confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=target_names,
            yticklabels=target_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title(f'Confusion Matrix - {dataset_name} Dataset')
plt.tight_layout()
plt.savefig(f"naive_bayes_{dataset_name.lower()}_confusion_matrix.png")

# Visualize probability distributions
# For simplicity, we'll just plot the first two features for each class
plt.figure(figsize=(12, 6))

# Create a DataFrame for easier plotting
df = pd.DataFrame(X, columns=feature_names)
df['target'] = y
df['target_name'] = [target_names[i] for i in y]

# Plot the first two features
plt.subplot(1, 2, 1)
sns.scatterplot(x=feature_names[0], y=feature_names[1], hue='target_name',
data=df)
plt.title(f'{feature_names[0]} vs {feature_names[1]}')

# Plot another set of features if available
if len(feature_names) > 3:
    plt.subplot(1, 2, 2)
    sns.scatterplot(x=feature_names[2], y=feature_names[3],
hue='target_name', data=df)
    plt.title(f'{feature_names[2]} vs {feature_names[3]}')

plt.tight_layout()
plt.savefig(f"naive_bayes_{dataset_name.lower()}_feature_distribution.png")

```

```
        return accuracy, report

def main():
    # Implement Naive Bayes on Iris dataset
    implement_naive_bayes("Iris")

    # Show the plots
    plt.show()

if __name__ == "__main__":
    main()
```

Visual Outputs



Iris Confusion Matrix



Iris Feature Distribution

Conclusion

Naive Bayes is a powerful classification algorithm despite its simplicity and "naive" assumption of feature independence. Through this experiment, we:

1. Successfully implemented Gaussian Naive Bayes for multi-class classification
2. Achieved high accuracy on both the Iris and Wine datasets
3. Visualized the model's performance using confusion matrices and feature distribution plots

The algorithm works particularly well for datasets with well-separated classes and when the independence assumption doesn't significantly impact the classification task. It's computationally efficient and requires relatively little training data to estimate the necessary parameters.

For more complex datasets with strong feature dependencies, other classification algorithms like Random Forests or Support Vector Machines might perform better. However, Naive Bayes remains a strong baseline algorithm and is particularly effective for text classification tasks like spam detection and sentiment analysis.